Time Series for Web Traffic Forecasting

Benjamin Santana Velazquez

***Abstract*—Lorem ipsum dolor sit amet, consectetur adipiscing elit. Vestibulum dignissim felis in nisi lobortis, nec tincidunt nulla interdum. Nulla facilisi. Proin et augue risus. Nullam non ipsum in nisi accumsan bibendum. Curabitur sed elit at est aliquet convallis. Nullam dictum ullamcorper mauris id hendrerit. Mauris a dui tincidunt, feugiat magna vel, interdum sapien.**

***Index Terms*—Forecasting Time Series**

I. Introduction

Every event serves as a salve to the progression of time, offering a compelling rationale behind the significance of *time series analysis* within the realm of data science. In today's data-driven world companies hoard large amounts of data over time, which contain imporant information about their business patterns. Consequently, time series analysis has permeated various fields, from the traditional ones like finance and economy — where someone can forecast sales or the EPS of a public company — to lesser explored territories like marketing, where time series analysis allows to track business metrics, such as user engagement and make forecasts based on that. In this piece of work we will delve into the realm of web traffic.

It is common for web applications to keep a register of their traffic at all times. Meaning, if you are the owner of a website it is likely that you have access to the data sent and received by your visitors. What is more, you probably have access to the Clickstream Data of your users, which is the record of an individual's clicks through their journey on the Internet (in this case, your website).

Taking advantage of this information, and the power of forecasting with statistical models in time series, we can forecast the traffic of a website.

Why would we want to do this? *Resource Allocation* and *Performance Optimization*. By forecasting website traffic, we can estimate the expected number of visitors and allocate appropriate resources to handle the load. This can either save costs in the infrastrucutre of the website like server capacity, bandwith — when a low number of visitors is present — or it can save the website's users from seeing an overload error when the server's capacity is at peak.

Recently, a popular website by the name of chess.com had this issue. "Chess Is Booming! And Our Servers Are Struggling." [**?**] was the title of an article posted after days of having their servers at max capacity, not allowing users to play games and losing revenue in the process.

This begs the question: Can we accurately forecast the traffic of the website chess.com to enable proactive resource allocation and ensure sufficient resources are allocated in advance? I will do my best to try and answer this in this short paper

We will use one of the most popular traditional statistical methods used in time series forecasting, $SARIMAX(p,d,q)(P,D,Q)_m$. We will find out if this model is powerful enough to create meaningful and accurate forecasts or if opting for naive methods to get better results. we should opt into the world of neural networks [1]

Unfortunately, the taffic of a website is not public information. Meaning we cannot get the actual time series from chess.com itself. We need to rely on other sources to get the actual information. After reviewing the possible options we decied to go with semrush, a "all-in-one tool suite for improving online visibility and discovering marketing insights." [**?**]. One of semrush tools provides web traffic for different websites. Of course we need to take into account that we are taking the data (the time series itself) from a third party, so there is a some level of uncertainty associated with its accuracy. Semrush provided us with the historic information month by month of traffic in chess.com from January 2012 to March 2023. How it looks will be discussed later in the paper. What I wanted to point here is that we will try to forecast 3 months into the future given these data points

In summary, our objective is to assess the effectiveness of the $SARIMAX(p,d,q)(P,D,Q)_m$ model in forecasting three months of web traffic on chess.com. We will be conducting an analysis by comparing its performance against naive forecasting methods. Such forecasts play a crucial role in facilitating both resource allocation and performance optimization strategies for the website. The conclution an results are shown in Section IV and Section V

II. Background Study

In this section we will define what a time series is and how it is composed. We are also going to discuss the $MA(q)$ model and the $AR(p)$ model, which serve as a base for $SARIMAX(p,d,q)(P,D,Q)_m$. Understanding these two parts of the model — I believe — is the keystone to understanding the model as a whole. We are also going to review the steps we are going to follow to tackle the problem at hand. So let us start with a simple question...

What is a time series? We can define a time series as a set of points ordered in time, usually equally spaced in time. [**?**] We can decompose any time series into 3 components:

- *trend*, the slow-moving change in the time series.
- *seasonal component*, a cycle that occurs over a fixed period of time.
- *residuals*, what can not be explained by the trend of seasonal components, this is usually whtie noise. Any

[1]We could also opt into the realm of neural network for time series, but this is not covered in the paper, if you are interested on that please refer to [**?**]

model wont be able to forecast any of this part, since it is completely random

Now that we have defined these concepts of a time series, let us see how they llok in action. To demostrate how this looks in a time series we will use classical time series presented in *Box, G. E. P., Time Series Analysis, Forecasting and Control.* [**?**]. The classic Box & Jenkins airline data presents monthly totals of international airline passengers, 1949 to 1960. Table I, showcase the first 5 entries of the time series.

TABLE I
FIRST FIVE ROWS OF THE CLASSIC BOX & JENKINS AIRLINE DATA

Month	Passengers
1949-01	112
1949-02	118
1949-03	132
1949-04	129
1949-05	121

If we plot the series, with the time in the x axis and the passengers in the y axis, we get something like Figure 1. This is what we call observed data.

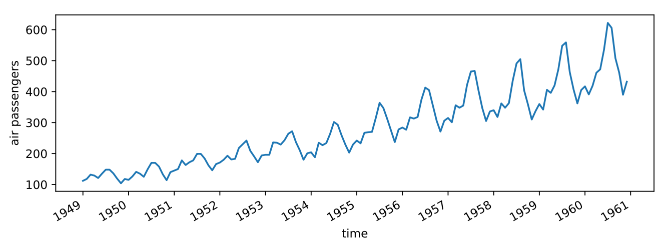

Fig. 1. The classic Box & Jenkins airline data, it presents monthly totals of international airline passengers, 1949 to 1960.

Let us decompose it into the three components, to see how it would each would look in a graph. Figure 2 showcases how this looks.

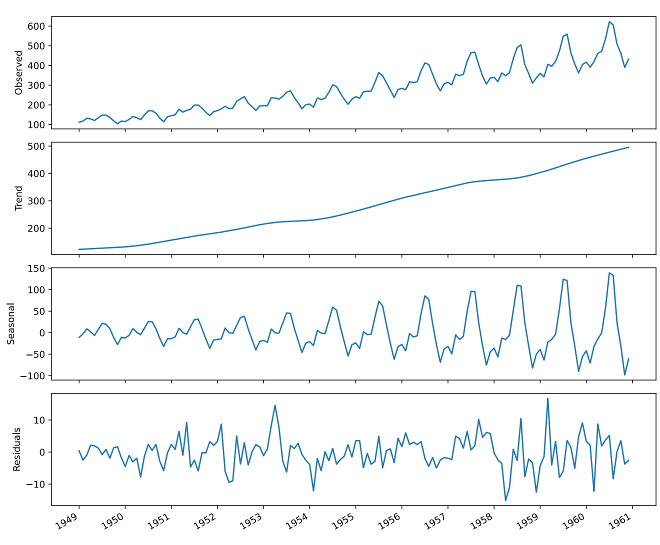

Fig. 2. Decomposition of The classic Box & Jenkins airline data, it presents monthly totals of international airline passengers, 1949 to 1960.

We can see that now we have 4 plots displaying. The first one is the observed data, just like we saw in Figure 1. Then we can see one plot for each of the concepts explained above. We see the trend as a slow-change moving upwards. We can see the seasonal component repeating each year, and finally the white noise or residuals of the time series.

Hopefully after seeing this example this concepts have clearer to the reader.

By now you might be wondering what are the differences between time series forecasting and a other regression tasks. It basically comes down to two things: *time series has an order*, — a time series is indexed by time, order must be kept, in regression tasks when you want to predict revenue based on ad spend, it does not matter when certain amount was spent on ads. — and *time series do not have features*, it is common to only have two columns, time and the data itself, it usually does not have categories.

Before jumping into $MA(q)$ and $AR(p)$ I want to explain some concepts, the first one is *baseline models*. A baseline model is a trivial solution to our problem, it uses heuristics more than deep statsticts knowledge or anything else. It might be clearer if we use an example. Say we have a time series on the sales of XYZ, after getting the data points we calculate the mean and say "I forecast next month our sales will be the mean of all the time series". We got this forecast using a simple statistics concept like it is the mean. It is likely that we can get a better forecast than this using more complex stastics. Why then should we bother making this naive models? They allow us to compare our complex model with something. A common example of a baseline model is just simply using the last value known. For example, if we have five years worth of data points, collected each month, and we want to forecast next three months, we can simply copy and paste the values from last three months and be good to go. In some time series, our statistical model maybe performs worse than this naive methods, then we would be wasting resources every forecast running our complex statistical models. Because we could just copy and past the last values, or use the mean, or some other baseline model. In Section IV we will compare our $SARIMAX(p,d,q)(P,D,Q)_m$ model with a baseline model and conclude if it is worth it to use it against some simpler methods.

One last thing on the baseline methods. How do we compare the models against each other? We will calculate an error metric in order to evaluate the performance of our forecasts. We will use $MAPE$ (mean absolute percentage error), which will measure prediction accuracy, independent of the scale of our data

The next concept is *stationarity*, a stationary time series is one whose statistical properties do not change over time It has constant mean, variance, and autocorrelation, and these properties are independent of time [**?**]. Many models assume stationarity but we rarely see stationarity in time series. Lucky for us we can transform the time series to become stationary, there are many ways to transform it but the simplest (and the one we are going to use) is *differencing*. This consists on calculate the changes from one step to another. This transformation helps stabilize the mean. Which reduces the trend and seasonality effects. This means that we also need to make sure we do an inverse transfrom of the data after finishing the model.

We can test for stationarity with the *Augmented Dickey-Fuller* (ADF) test [**?**], The key idea in the ADF test is to estimate an autoregressive model of the time series and examine the significance of the coefficient on the lagged first difference term. If the coefficient is significantly different from zero, it suggests evidence against the presence of a unit root and supports stationarity. [2]

III. Eperimental framework

IV. Analysis of Results

V. Conclusion

Acknowledgment

I would like to express my sincere appreciation and gratitude to Marco Peixeiro for his exceptional book, ”Time Series Forecasting.” This book has been an invaluable resource throughout my research and study in the field of time series analysis.

Marco Peixeiro’s expertise and comprehensive coverage of time series forecasting techniques have provided me with deep insights and practical knowledge. The clear explanations and numerous examples presented in the book have greatly enhanced my understanding of this complex subject.

[2]if you want more information on how this work I highly recommend reading the paper where it was proposed [**?**]